Fundamentals of Data Engineering

Trainer: Nilesh Ghule

Big Data
5 que

Quiz 1 → 15th
Quiz 2 → 16th
} 10 mcq.

D.E

module end quiz } 20 mcq.

1. Databases
2. Big Data
3. Cloud Computing
4. ...

Pre-requisites

1. Computer fundamentals
2. Operating Systems
3. Networking
4. Programming : Python, Java, ...

Introduction

- **Big Data Fundamentals**
  - Evolution of Data Engineering | V's: Volume, Velocity, Variety, Veracity, Value
- **Databases**
  - RDBMS - ACID, SQL (basic concept only) I NoSQL - BASE, CAP theorem
- **Data warehouse - OLAP vs OLTP**
  - Data cleansing, Data transformations and Data modelling I Data warehouse vs Data mart
- **Data Engineering Life Cycle**
  - Source → Ingestion → Storage → Transformation → Serving
  - Ingestion: ETL vs ELT
  - Storage: Distributed storage, Storage services I Processing: Batch vs Stream
- **Cloud computing fundamentals**
  - Virtualization, Scaling, Elasticity, Cloud service models, Vendors
- **Big Data Technologies**
  - Frameworks: Hadoop, Hive, Spark, Kafka
  - Applications and Job profiles.

Data Engineering at a Glance

Database & Warehouse

- File IO - Data
- 1970: E.F. Codd
  - → Relational DBMS
- 1980: RDBMS popular
- 1991: Data Warehouse
  - ↳ Data analysis.

SQL

Internet & DotCom

- 1993: WWW
  - business on internet.
- 1995: Java popular

NoSQL Database

- 1998: Carlo Strozzi
- High performance, Huge Data, ...

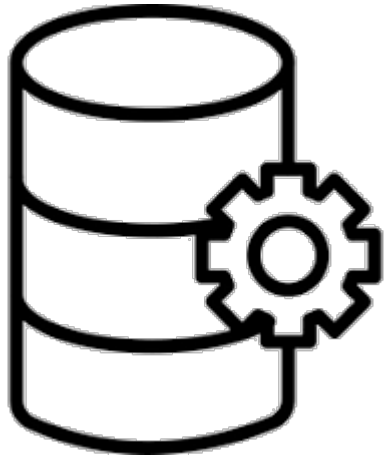

MPP & Big Data Tech

- multi-core: 2000s
- GPU for data process: 2005
- Grid Computing - Super Computer
- Distributed Computing Storage

Cloud Computing

- Data Centers
  - ↳ AWS, GCP, Azure, Smart
- Pay per use.

Big Data characteristics

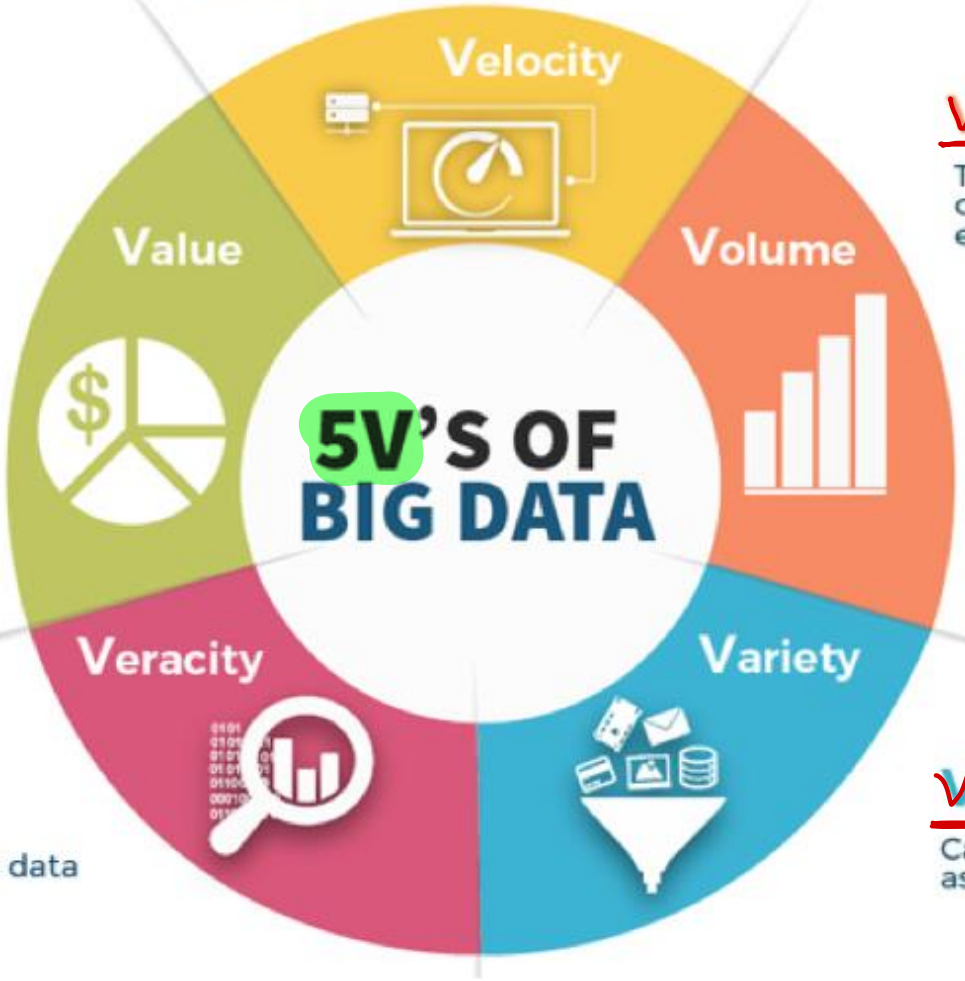

Velocity
Speed at which data is emanating and changes are occurring between the diverse data sets

Volume 100s TB+
This refers to the sheer volume of data being generated every second.

Value
Having access to big data is all well and good but that's only useful if we can turn it into a value.

Veracity
Data reliability and trust.
Verifying and validating the data

Variety
Can use structured as well as unstructured data.

Types of Data

files have fixed formats but contents cannot be understood programmatic (traditional).

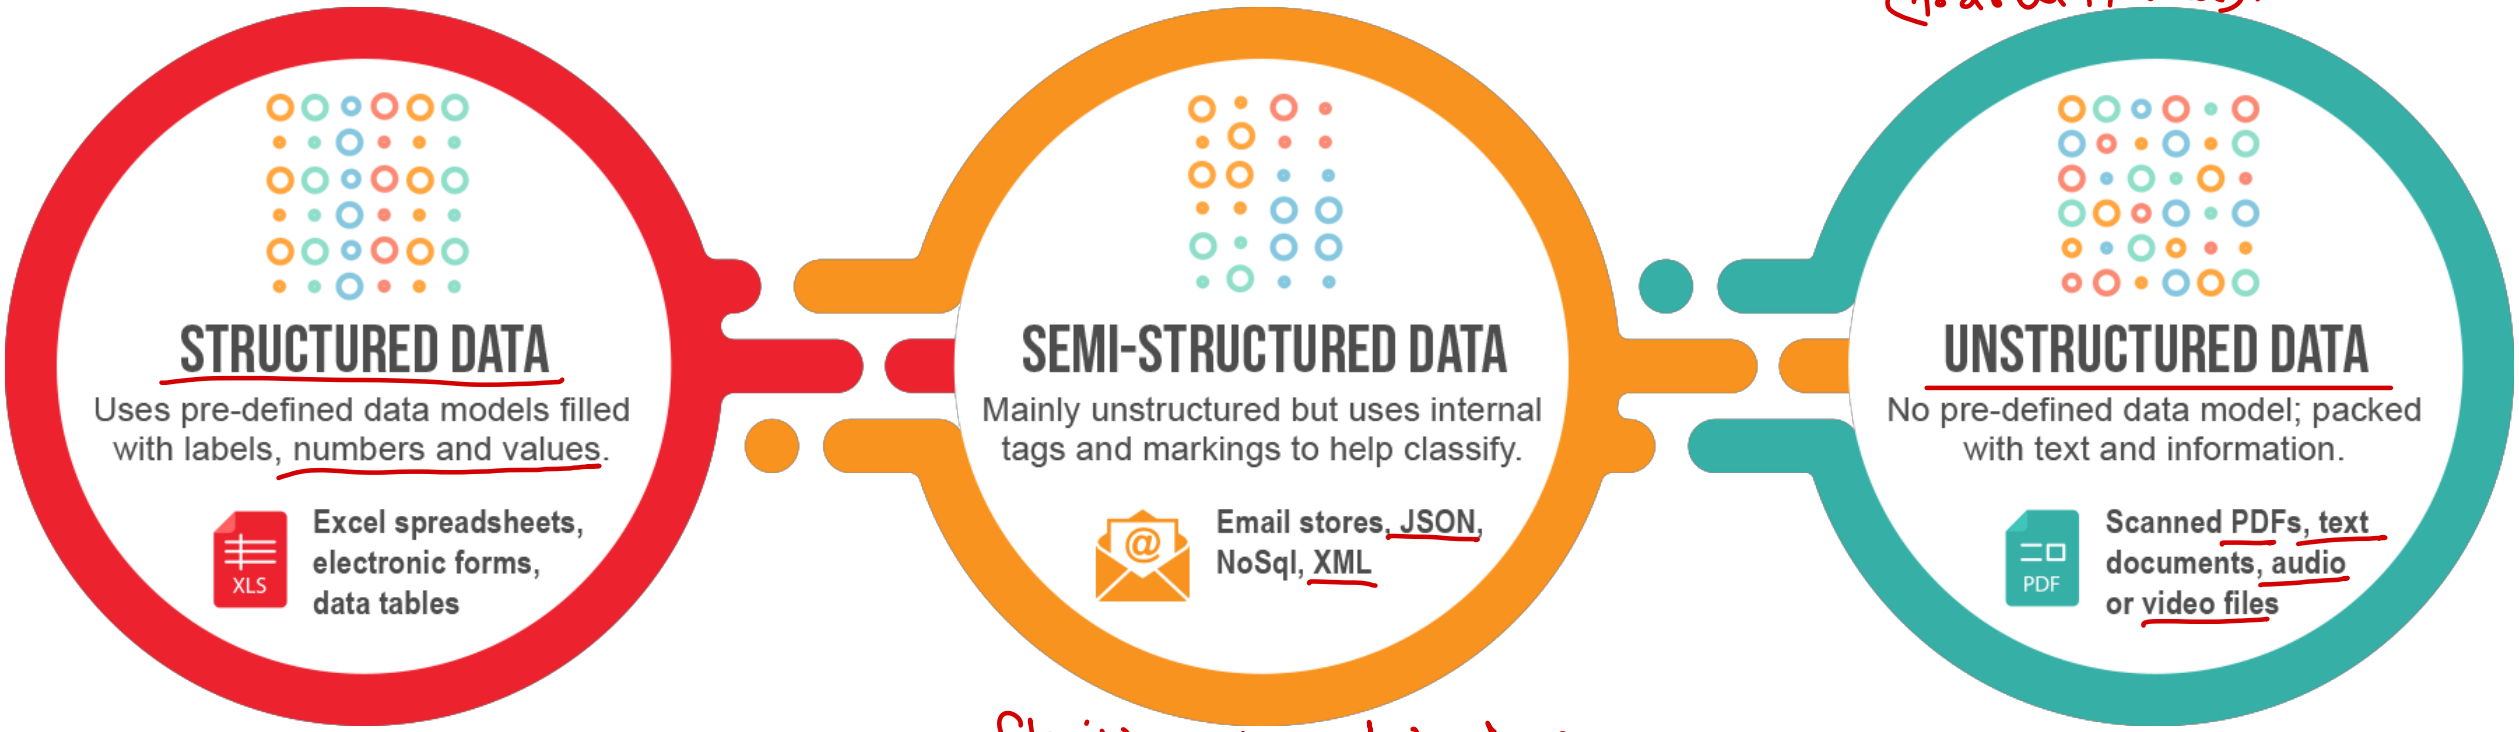

flexible schema/structure.

RDBMS

DBMS : CRUD ops → Create/Insert, Retrieve/Select, Update/change, Delete
1970 : Codd rules → RDBMS

- Every enterprise application need to manage data.
- RDBMS is relational DBMS than manages structured data.
- Data is organized into tables, rows and columns. Tables are related to each other.
- All enterprise RDBMS follow server-client architecture, have built-in relational capabilities, fully ACID transactions, based on Codd's rules.
- DB2, Oracle, MS-SQL, MySQL, Postgre-SQL, MS-Access, SQLite, etc.

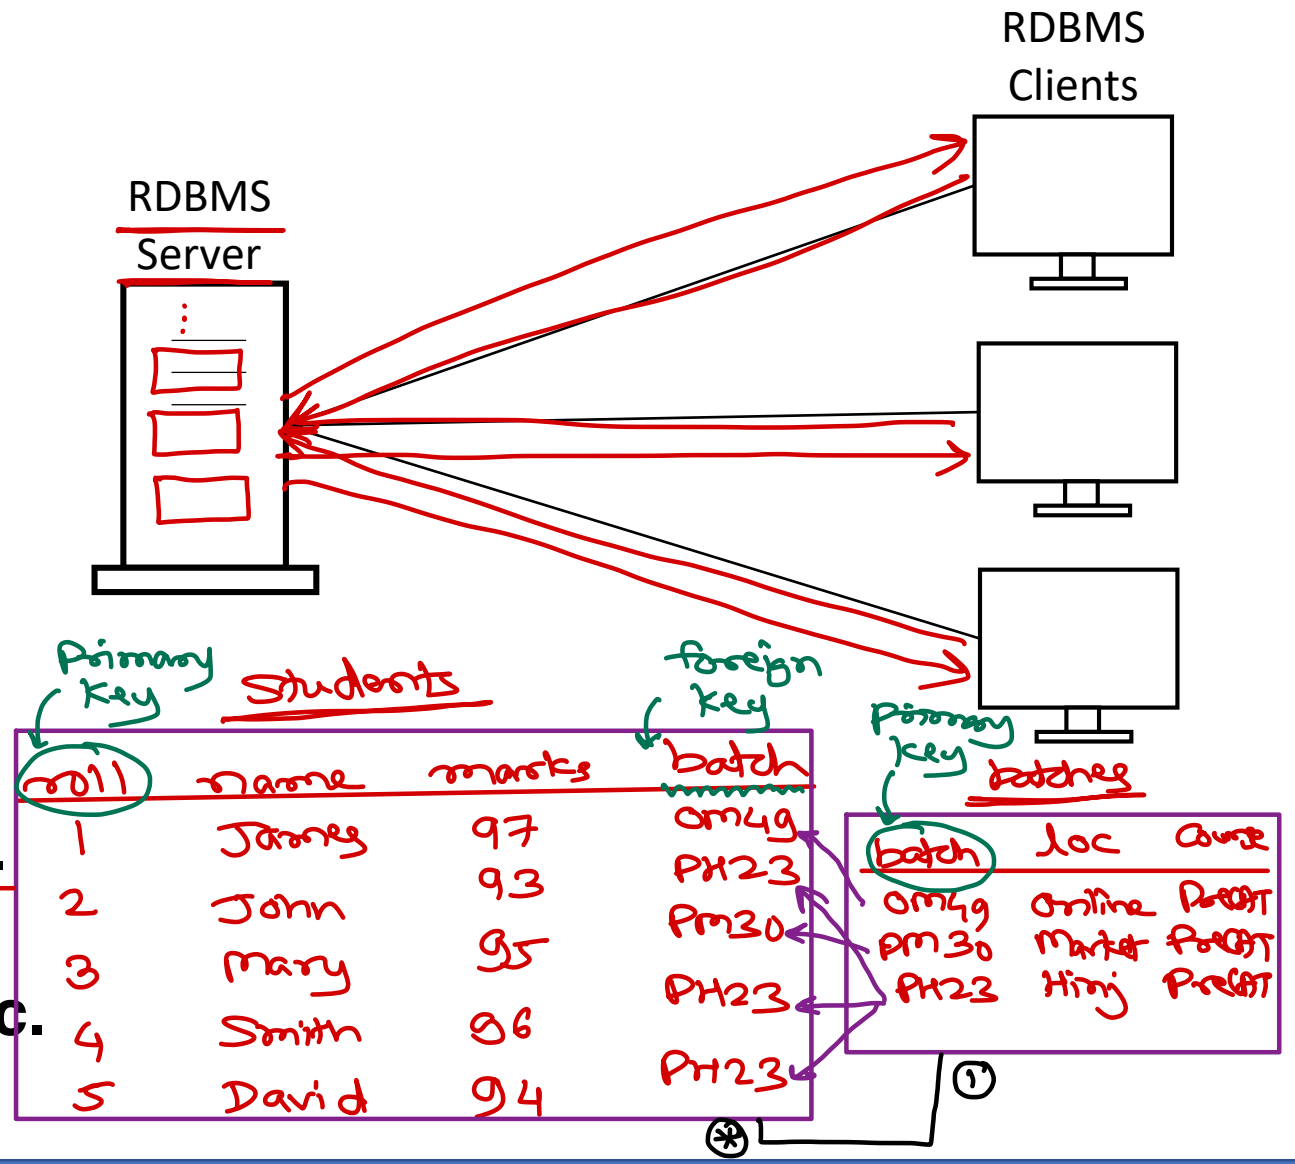

SQL – Structured Query language

- RDBMS data is processed with SQL queries.
- ANSI standardised in 1986. / ISO std in 1987.
- Five major categories:
  - DDL: Data Definition Language e.g. CREATE, ALTER, DROP, RENAME. ← related to structure / schema of data.
    - CREATE TABLE people(id INT, name CHAR(40), birth DATE);
  - DML: Data Manipulation Language e.g. INSERT, UPDATE, DELETE.
    - INSERT INTO people VALUES(1, 'Nilesh', '1983-09-28');
    - UPDATE people SET name='NILESH' WHERE id=1;
    - DELETE FROM people WHERE id=1;
  - DQL: Data Query Language e.g. SELECT.
    - SELECT * FROM people;
  - DCL: Data Control Language e.g. CREATE USER, GRANT, REVOKE.
  - TCL: Transaction Control Language e.g. SAVEPOINT, COMMIT, ROLLBACK.

RDBMS — ACID Transactions

accounts

id	type	balance
1	Saving	10000
2	Saving	2000
3	Current	50000
4	Saving	20000

3000

accounts

id	type	balance
1	Saving	~~10000~~ 7000 ← UPDATE
2	Saving	~~2000~~ 5000 ← UPDATE
3	Current	50000
4	Saving	20000

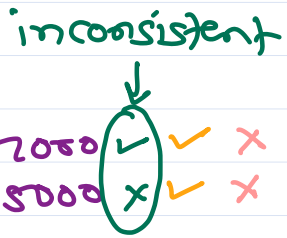

1. START TRASACTION;
2. UPDATE accounts SET bal=7000 WHERE id=1;
3. UPDATE accounts SET bal=5000 WHERE id=2;
4. COMMIT; or ROLLBACK;

Tx is set of DML queries that is executed as a single unit.
Either all queries in tx should be completed → COMMIT
OR
all queries in tx should be discarded. → ROLLBACK

1. ATOMIC → All queries success or fail.
2. CONSISTENT → Same data visible to all clients.
3. ISOLATED → Multiple tx will execute independently.
4. DURABLE → Final results must be stored on disk.

Scalability

Scalability
- Number of clients
- Amount of data

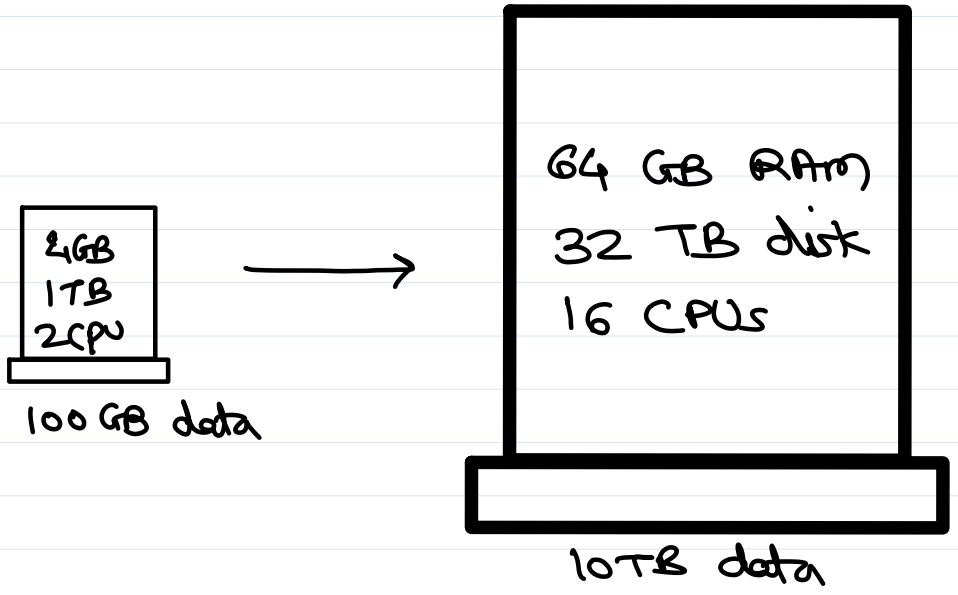

Vertical scaling
OR
Up-Scaling

RDBMS scaling

Distributed systems

Horizontal scaling
OR
Out-Scaling

NoSQL scalability

cluster = set of computers for dedicated task in a network

NoSQL Databases

1998 - Carlo Strozzi
his own db without SQL

NoSQL = Anti-SQL.

Last.fm - international conference
2009 - invitation #nosql
- nosql become popular.

- Stands for Not Only SQL
- Manages structured and semi-structured data.
- Prioritizes high performance, high availability and scalability
- Designed for Horizontal scaling. Reliable, fault tolerant, Better performance/Speed.
- No declarative query language → different languages for diff dbs.
- Uses: Huge data (TBs), Many Read/Write ops, Scalable, Flexible schema.
- Don't use if: Need high consistency, Multiple relations
- BASE transactions and Based on CAP Theorem

CAP Theorem – a.k.a. Brewer's Theorm

- **Consistency - Data is consistent after operation. After an update operation, all clients see the same data.**
- **Availability - System is always on (i.e. service guarantee), no downtime.** 24x7
- **Partition Tolerance - System continues to function even the communication among the servers is unreliable.**

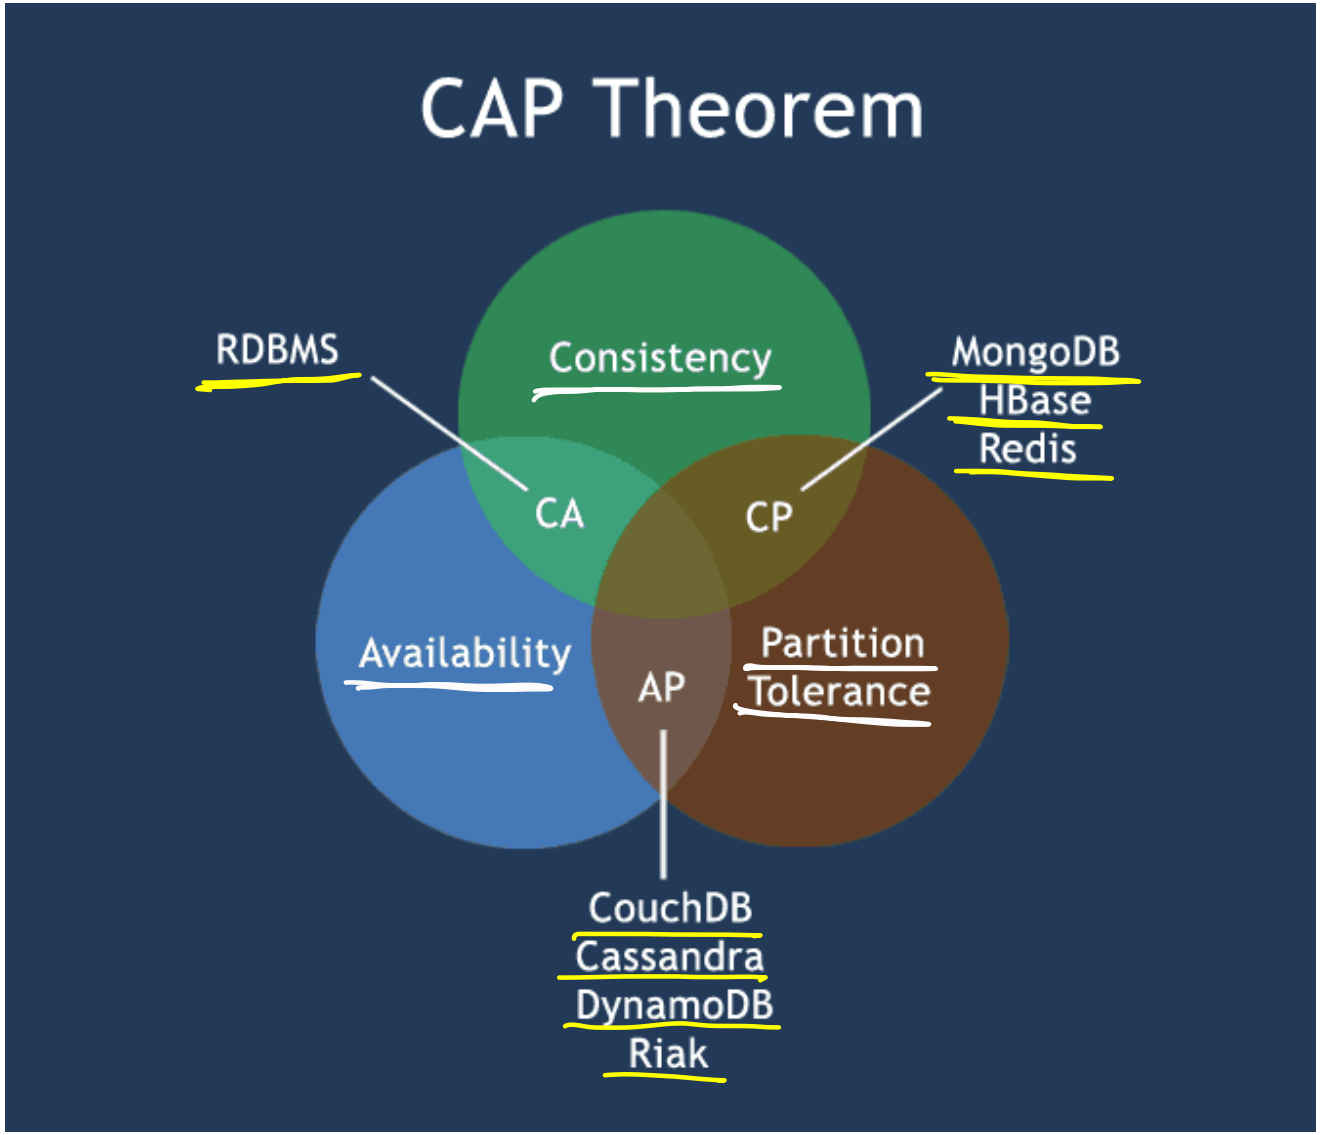

NoSQL — Partition Tolerance and BASE Transactions

Partition Tolerance

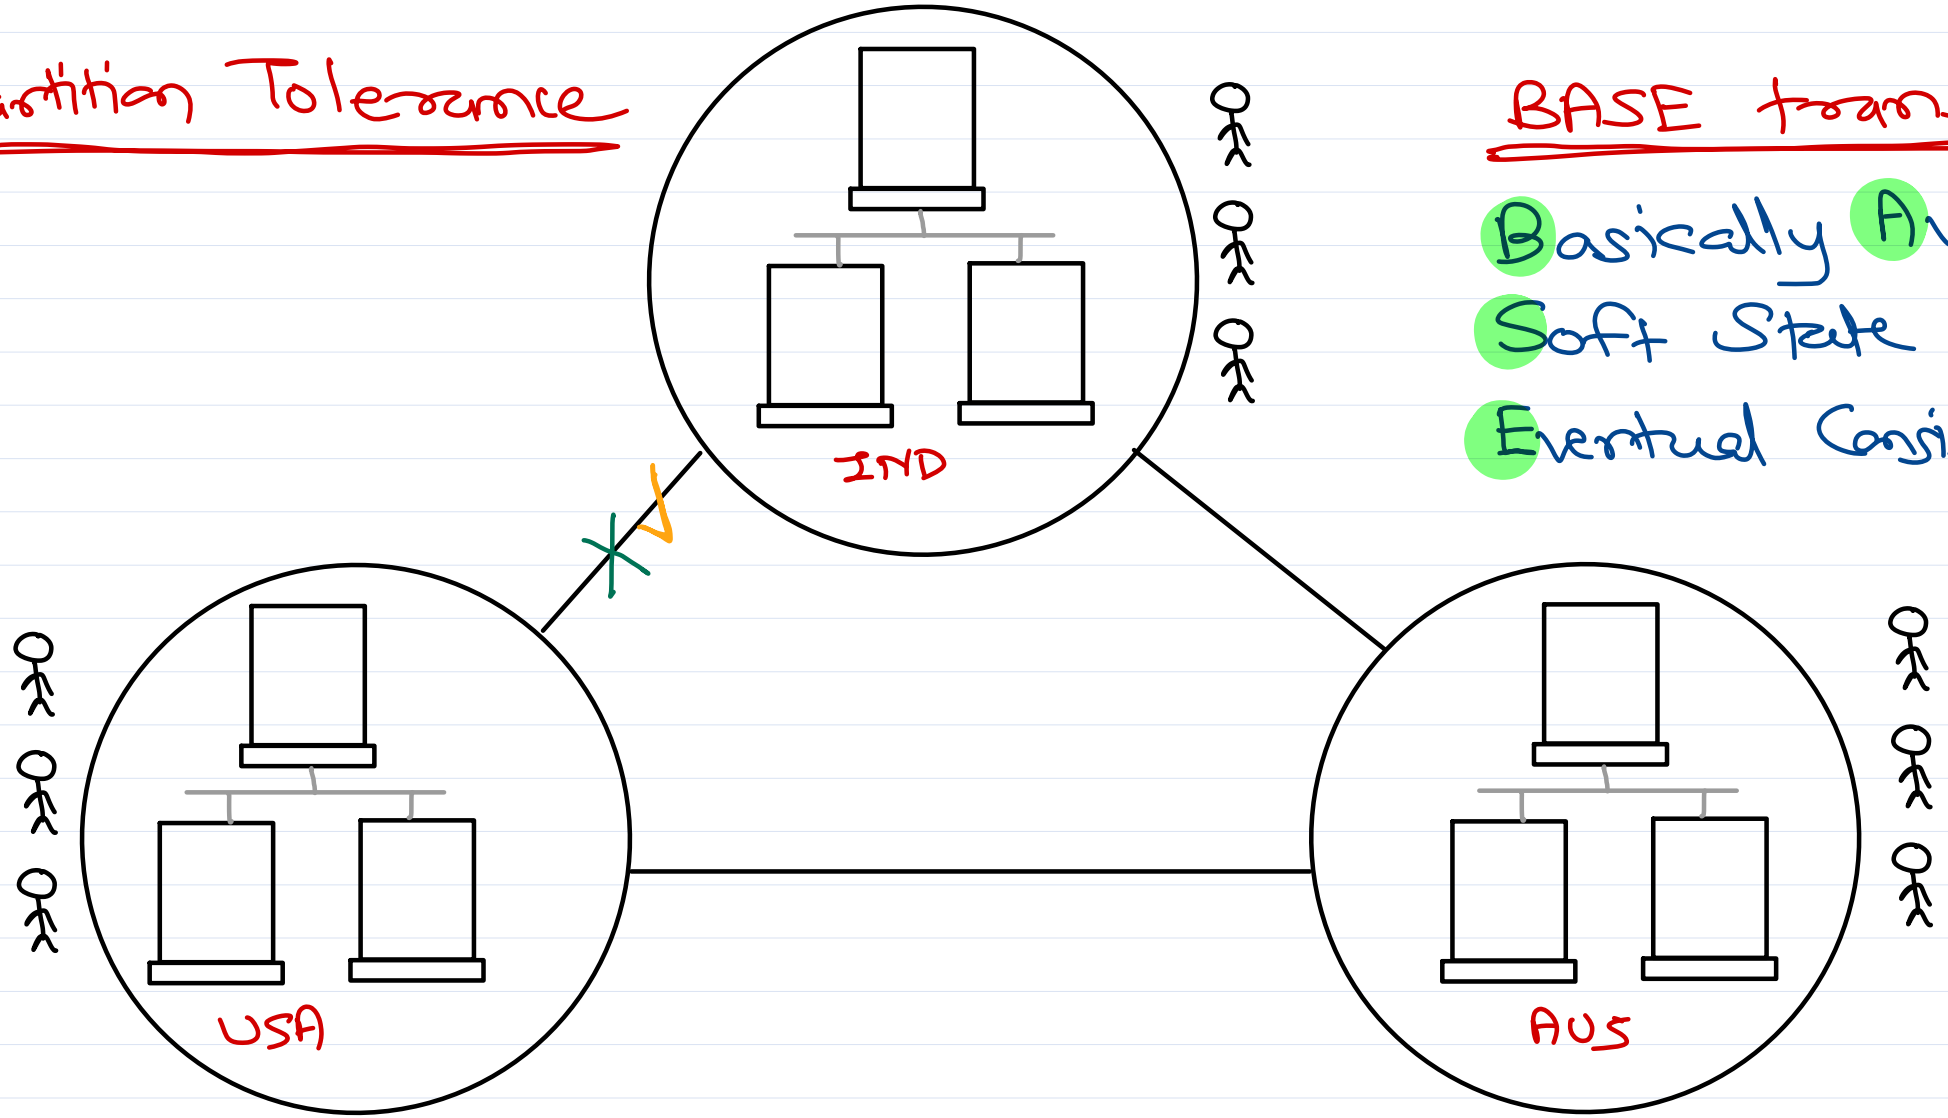

BASE transactions

- **B**asically **A**vailable
- **S**oft State
- **E**ventual Consistency

Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>